

Reinforcement Learning

Exercise 4

Mathias Niepert, Vinh Tong

April 28, 2023

1 Monte Carlo Methods vs Dynamic Programming (3P)

- a) What are advantages of Monte Carlo methods over dynamic programming? Mention at least two. (2P)
- b) Give an example environment where you would use a Monte Carlo method to learn the value function rather than using dynamic programming. Explain why. (1P)

2 Monte Carlo ES for blackjack (6P)

In this exercise we use the blackjack environment from gym (https://www.gymlibrary.ml/environments/toy_text/blackjack/). The code template can be found on Ilias in `ex04-mc/ex04-mc.py`. It has been tested with gym version 0.18.0 (but should also be stable with version 0.18.3).

- a) Consider the version of blackjack introduced in the lecture (Example 5.1 from Sutton and Barto). Implement first-visit Monte Carlo prediction (lecture 4 slide 13) for the given policy: stick if sum ≥ 20 , else hit. Try to reproduce the figures on slide 16. (3P)
- b) Implement Monte Carlo ES (slide 23) and obtain the optimal policy and state-value function for blackjack. Output the policy every 100,000 iterations (e.g. as 2 tables, one with usable ace and one without usable ace). We recommend an optimistic initialisation of Q to improve results. Let it run for at least 500,000 iterations. Letting it run until convergence might take a very long time, so intermediate results are okay! (3P)